
CrackMeBench: Binary Reverse Engineering for Agents

Isaac David
University College London

Arthur Gervais
University College London

Abstract

Benchmarks for coding agents increasingly measure source-level software repair, and cybersecurity benchmarks increasingly measure broad capture-the-flag performance. Classical binary reverse engineering remains less precisely specified: given only an executable, can an agent recover validation logic and produce an input, serial, artifact, or key generator accepted by the program? We introduce CrackMeBench, a benchmark for evaluating language-model agents on educational CrackMe-style reverse-engineering tasks. CrackMeBench focuses on deterministic binary validation problems with executable oracles, symbol-poor binaries, explicit local tool access, and externally scored submissions rather than free-form explanations. The v0 benchmark combines eight public calibration CrackMes with twelve generated main-score tasks built from seeded C, Rust, and Go templates, and agents run through an equal shell interface in a no-network Linux Docker sandbox with standard reverse-engineering tools. In a three-model evaluation with a five-minute budget and three scored submissions per task, pass@3 on the generated split is 11/12 tasks (92%) for GPT-5.5, 7/12 (58%) for Claude Opus 4.7, and 5/12 (42%) for Kimi K2. The harder generated half separates the models more sharply, with pass@3 of 5/6, 2/6, and 1/6, respectively; on the eight-task public calibration split, pass@3 is 3/8, 2/8, and 1/8. CrackMeBench records pass@1 and pass@3, scored submissions, wall-clock time, command traces, tool categories, provider-reported token usage, estimated cost, and qualitative failure labels, providing a reproducible testbed for measuring progress from source-code reasoning toward autonomous binary analysis while restricting scope to educational, purpose-built programs.

1 Introduction

Language-model agents are increasingly evaluated as programmers, security analysts, and autonomous tool users. Source-code benchmarks such as HumanEval, MBPP, APPS, AlphaCode, DS-1000, LiveCodeBench, and SWE-bench measure whether models can synthesize functions, solve programming problems, generate data-science code, avoid contamination, or repair repositories from issue descriptions [Chen et al., 2021, Austin et al., 2021, Hendrycks et al., 2021, Li et al., 2022, Lai et al., 2023, Jain et al., 2025, Jimenez et al., 2024]. Cybersecurity benchmarks such as CyberSecEval, NYU CTF Bench, and Cybench instead place agents in adversarial puzzle environments that span secure coding, web exploitation, binary exploitation, forensics, cryptography, reverse engineering, and miscellaneous tasks [Bhatt et al., 2024, Shao et al., 2024, Zhang et al., 2025]. These benchmark families have driven rapid progress, but they leave a narrower and older reverse-engineering question under-specified: when the main artifact is a compiled executable rather than a source repository or a broad CTF service, can an agent recover validation logic and produce an input, serial, artifact, or key generator that the executable accepts?

This question matters because binary reverse engineering stresses capabilities that are only partially exercised by source-level coding tasks. An agent must triage a file format, identify architecture and linkage, reason over stripped or symbol-poor code, decide when decompiler output is trustworthy,

formulate and test hypotheses dynamically, and often invert transformations or constraints that were not written for readability. Unlike source repair, the agent cannot rely on repository tests or semantic names. Unlike broad CTF evaluation, the objective can be made narrow and deterministic: produce the password, serial, input file, or key generator that the original executable accepts.

CrackMe programs are a natural educational substrate for this problem. CrackMes are purpose-built reverse-engineering exercises that commonly involve password checks, serial validation, anti-debugging, packing, or keygens [Crackmes.one, 2026]. They are small enough for controlled benchmarking, but rich enough to expose failure modes that are common in real reverse-engineering workflows: over-trusting decompilers, missing input/output conventions, confusing encoded data with final strings, hardcoding one example instead of recovering a keygen, and failing to combine static and dynamic evidence.

We propose CrackMeBench, an Evaluations & Datasets benchmark centered on executable-oracle binary validation. In each task, an agent receives a binary, public metadata, an explicit list of installed tools, a time budget, and a submission schema. The agent interacts with a local shell in a no-network Docker sandbox. It may inspect, disassemble, trace, emulate, script, or symbolically execute the binary. It must eventually submit a JSON object naming a password, artifact path, or keygen script. Scoring is performed by a private oracle outside the agent-visible workspace.

The executable-oracle framing is important. Written explanations are useful for qualitative analysis, but they are not the primary score: a candidate password must run successfully, an artifact must be accepted by the binary, and a keygen must generalize to hidden usernames. This sharply reduces ambiguity in evaluation and lets failures be assigned to operational categories such as failed triage, found-check-but-not-inverse, symbolic setup failure, or wrong submission format.

CrackMeBench is intentionally not a benchmark for malware analysis, commercial software cracking, DRM circumvention, or intrusion. The v0 scope is educational CrackMe-style validation logic, run locally with no network and no external targets. This restriction is both ethical and methodological: it lets the benchmark isolate reverse-engineering behavior without rewarding offensive use.

Our contributions are:

- We formulate CrackMe-style binary validation as an executable-oracle agent evaluation problem, where success requires a password, artifact, or key generator accepted by the original program rather than a plausible prose explanation.
- We introduce a reproducible benchmark and harness that combine public calibration CrackMes with generated main-score tasks, explicit tool manifests, no-network Docker execution, hidden oracle checks, command traces, token accounting, and structured failure labels.
- We report a three-model empirical evaluation showing that GPT-5.5 ranks first on generated CrackMes, Claude Opus 4.7 ranks second, Kimi K2 ranks third, and all evaluated models remain weak on public calibration CrackMes under the v0 protocol.

2 Benchmark Design

CrackMeBench is designed to be narrow enough for controlled measurement and broad enough to expose distinct reverse-engineering workflows. This section defines the task abstraction, the v0 task inventory, the generated/public split, and the oracle interface.

2.1 Design Goals

The benchmark is organized around four design goals. First, success should be externally verifiable. The final answer is not a prose explanation of a suspected password check, but a candidate that is executed against the original program or private oracle. Second, the task interface should resemble the common working conditions of binary reverse engineering: symbol-poor executable input, local tools, scratch scripts, and iterative dynamic tests. Third, tool access should be explicit. The agent should not infer the availability of Ghidra, angr, radare2, or tracing tools from model priors; the exact manifest is generated from the container and included in the prompt. Fourth, the benchmark should separate public calibration from main-score evaluation. Public CrackMes make the benchmark easier to inspect and compare with recognizable exercises, but they are not clean private evaluation data because comments and writeups can leak solutions.

ID	Split	Size (B)	LoC	Submission	Mechanism
P01	public	14,736	–	password	public fixed-password Linux smoke test
P02	public	828,056	–	keygen	public deterministic name-to-serial keygen
P03	public	320,192	–	keygen	public Rust keygen with runtime/decompiler noise
P04	public	180,416	–	password	public assembler and control-flow password task
P05	public	29,128	–	artifact	public VM-style artifact and bytecode task
P06	public	166,320	–	artifact	public crypto/artifact validation task
P07	public	21,136	–	password	public C++ string-manipulation password task
P08	public	15,480	–	password	public XOR-combination argv password task
S01	generated	14,416	17	password	argv password compared against stripped rodata literal
S02	generated	14,416	29	password	encoded byte array inverted through XOR and rotate
S03	generated	14,416	27	password	fixed input recovered from byte-wise affine checks
S04	generated	14,424	33	keygen	username-to-serial algorithm using 32-bit mixing
S05	generated	14,416	45	password	permuted byte checks behind opaque predicates
S06	generated	14,424	41	password	password hash guarded by ptrace and timing checks
S07	generated	14,432	42	password	runtime-decrypted validator records
S08	generated	14,432	54	artifact	toy RSA-style artifact validation
S09	generated	14,416	39	password	tiny bytecode interpreter with masked character tests
S10	generated	14,416	41	password	rolling checksum with table lookups and checkpoints
S11	generated	383,384	28	keygen	Rust username-to-serial validator with runtime noise
S12	generated	1,245,319	37	keygen	Go username-to-serial validator with runtime noise

Table 1: CrackMeBench v0 task inventory. Public tasks are calibration tasks bundled as sanitized binaries; generated tasks are the main-score tasks built from seeded templates. Binary size is reported in exact bytes because the small generated Linux ELF’s are dominated by fixed dynamic-linking and ELF-header overhead. LoC is reported for released generated sources and omitted for public binary-only tasks.

2.2 Task Abstraction

Each task consists of a compiled executable, public metadata, public handout files if needed, a submission schema, and a private oracle. The agent-visible metadata specifies the task identifier, platform, binary path, objective, allowed submission types, wall-clock budget, scored-submission budget, and any runtime note needed for reproducible execution. The agent does not receive source code, hidden oracle data, public web comments, or writeups.

The v0 benchmark supports three submission forms:

```
password: {"type": "password", "value": "candidate"},
artifact: {"type": "artifact", "path": "/work/candidate.bin"},
keygen: {"type": "keygen", "path": "/work/solve.py"}.
```

For keygen tasks, the submitted script is invoked as `/work/solve.py USERNAME` and must print only the serial or key. The oracle samples hidden usernames and checks them against the original validation routine, so a hardcoded example is insufficient.

2.3 Task Inventory

The v0 benchmark contains twenty tasks: twelve generated main-score tasks and eight public calibration tasks. Table 1 summarizes the inventory. The generated split is divided into six core tasks and six harder tasks. The core tasks probe literal recovery, reversible encodings, arithmetic constraints, simple keygens, benign anti-debugging, and a mini VM. The harder tasks add opaque predicates, runtime-decrypted validators, toy cryptographic artifacts, checksum/table reasoning, and Rust/Go runtime noise. Most generated C ELF’s are close in size because fixed ELF headers, section layout, and dynamic-linking overhead dominate a few dozen lines of validation code; the Rust and Go generated tasks are substantially larger because their runtimes are linked into stripped binaries. We report exact byte counts rather than rounded kilobytes to make this clear. The public tasks provide external validity and reviewer-recognizable CrackMe workloads, but they are reported as calibration because public CrackMe pages can contain comments, uploaded solutions, or external writeups.

Generated main-score tasks. `S01_literal_rodata` is a stripped ELF that compares argv against a literal string in read-only data. It is a smoke test for file triage, strings, and xrefs, and should be

solvable without heavy tooling. `S02_encoded_string` stores an encoded byte array and checks the inverse of XOR plus rotate operations; it tests whether an agent can recognize that an apparent byte array is a transform, not the final password. `S03_linear_constraints` validates a fixed-length input through byte-wise affine constraints suitable for either manual inversion or solver scripting. `S04_bitvector_keygen` computes a username-dependent serial through a seeded 32-bit mixer and bit rotation, and therefore tests whether the agent can recover an algorithm and write a reusable script. `S06_anti_debug_timing` uses benign `ptrace` and timing checks before a hash comparison; it is not intended to be evasive malware, but it exercises the common need to reason about dynamic-analysis side effects. `S09_mini_vm` embeds a tiny bytecode interpreter that checks input length and masked character constraints, stressing indirect semantics without requiring a large VM.

Harder generated tasks. `S05_opaque_branch_maze` distributes byte checks behind opaque predicates and misleading control flow, forcing agents to distinguish semantic constraints from dead branches. `S07_packed_xor_loader` keeps validator records encrypted until runtime, so the direct strings and constants seen by static triage are not sufficient. `S08_toy_crypto_artifact` asks for a binary artifact accepted by a deliberately breakable small-RSA-style check rather than a command-line password. `S10_symbolic_resistant_checksum` combines rolling state, table lookups, and intermediate checkpoints that are awkward for naive symbolic execution but amenable to targeted scripting. `S11_rust_serial` and `S12_go_serial` implement username-dependent serial validation in Rust and Go, respectively, adding runtime and string-reference noise while preserving the same keygen oracle interface as `S04`.

Public calibration tasks. The public split contains Toronto’s *c*, ryndrka’s *Enigma CrackMe v1.0*, TheSwedishLord’s *SirCrackaLot v2*, tdaron’s *Use your brain*, Ben_Lolo’s *FlipVM*, victormeloasm’s *AI – Almost Impossible*, LoZ’s *Password Login System*, and Yandere’s *XOR*. These tasks cover simple password recovery, keygen recovery, Rust reverse engineering, assembler/control-flow analysis, VM artifact recovery, crypto/artifact validation, C++ string manipulation, and argv-based XOR reconstruction. They are bundled only as sanitized challenge artifacts with SHA-256 metadata and Apache-2.0 redistribution metadata; comments, web pages, source bundles, and solution files are not mounted for agents.

2.4 Generated Split and Release Policy

The generated tasks are built from seeded templates. A split-specific base seed controls passwords, constants, hidden usernames, bytecode tables, and artifact parameters. The released generator writes `metadata.yaml`, host-side oracle data, task source, and a README for each task; a Dockerized build step compiles C sources as Linux x86-64 ELF binaries with optimization and stripped symbols, and compiles the Rust and Go keygen tasks with their standard toolchains before stripping where supported. This makes the generated tasks reproducible while keeping the agent-facing artifact close to a CrackMe binary rather than a source exercise.

The artifact release includes generators, dev seeds, generated test binaries, generated test sources, metadata, Docker build scripts, oracle interfaces, and reviewer-accessible oracle data for the reported run. This design separates reproducibility from solution leakage: reviewers can rebuild and inspect the benchmark, but agents evaluated under the official protocol see only the compiled binary and public metadata. Future leaderboards can rotate or hide generated test seeds while preserving the same generator and oracle interface. Public CrackMes are handled differently: because they are externally authored assets, the repository records source URLs, SHA-256 hashes, acquisition notes, and license metadata, and mounts only sanitized challenge artifacts during evaluation.

2.5 Executable Oracles

Oracles run outside the agent-visible workspace against pristine task files. For fixed-password tasks, the oracle executes the original binary with the candidate input and checks exit status and accepted output. For keygen tasks, the oracle samples hidden usernames and validates each emitted serial. For artifact tasks, the oracle copies the candidate artifact into an isolated run directory and executes the relevant binary or VM invocation. We use *submission* to mean one candidate JSON file queued through `/harness/submit` and scored by the oracle; exploratory shell commands, local tests, and

Host orchestrator loads task metadata, model config, and <code>tools_manifest.yaml</code>; starts Docker; injects the exact tool list into the prompt; sends model commands to the container; logs outputs, timing, token usage, and tool categories; scores queued submissions through private oracles. Docker sandbox runs with no network, <code>/task:ro</code>, <code>/work:rw</code>, and <code>/harness:ro</code>; exposes standard reverse-engineering tools; queues submissions but does not contain oracle secrets.

Figure 1: End-to-end CrackMeBench execution pipeline. The model never receives direct filesystem access outside the sandbox, and the oracle never runs inside the agent-visible mount.

scripts are not submissions. Invalid JSON, missing files, non-executable keygens, and exceeded submission budgets are scored as protocol failures.

3 Harness and Evaluation Protocol

The harness and evaluation protocol jointly define what an agent can see, which tools it can use, how shell actions are executed, and how submissions are scored. The harness is designed to make tool access explicit, reproducible, and equal across models. Model API calls run on the host. All shell commands requested by the model execute inside a Linux x86-64 Docker container with network disabled. The agent-visible filesystem has three mounts:

- `/task:ro`: sanitized task metadata, the challenge binary, hashes, and public handout files;
- `/work:rw`: the agent scratch directory, command outputs, scripts, and submissions;
- `/harness:ro`: public helper commands `list_tools`, `help`, `submit`, and the tool manifest.

Private oracle files, task sources, model credentials, and host logs are not mounted inside the container.

End-to-end flow. At run start, the host loads environment variables, task metadata, model configuration, and the public tool manifest. It starts the Docker container with `/task`, `/work`, and `/harness`; executes `/harness/list_tools --json` inside that same container; injects the returned tool list into the model prompt; and then alternates between model calls and shell-command execution. When the model writes `/work/submission.json` and calls `/harness/submit`, the public submit wrapper queues the JSON in `/work/submissions`. The host detects the queued file and invokes the private oracle outside the container against pristine task files. This design keeps the shell interface simple while preventing the agent from reading hidden usernames, valid passwords, oracle constants, or model credentials.

3.1 Tool Disclosure

At the beginning of every run, the host executes `/harness/list_tools --json` inside the container and injects that exact JSON into the task prompt. This prevents ambiguity about which tools are installed and avoids relying on model priors. The v0 image includes static triage utilities (`file`, `strings`, `readelf`, `objdump`, `llvm-objdump`, `nm`, `xxd`), disassembly/decompilation tools (`radare2/r2pipe`, `rabin2`, `rasm2`, headless Ghidra wrappers), debugging and tracing tools (`gdb`, `strace`, `ltrace`, `qemu-x86_64`), Python reverse-engineering packages (`angr`, `Z3`, `claripy`, `capstone`, `unicorn`, `keystone`, `pyelftools`, `pwntools`), and build/scripting tools including `GCC/Clang`, `Rust`, `Go`, `CMake`, `Ninja`, and standard Unix utilities [National Security Agency, 2026, `radare2` Project, 2026, Shoshitaishvili et al., 2016, Luk et al., 2005, Nethercote and Seward, 2007]. Command output is stripped of ANSI/control sequences before being returned to the model, reducing accidental context blow-up from interactive tools.

3.2 Prompt and Command Protocol

All models receive the same system prompt: they are told that they are solving an educational CrackMe task in a no-network Linux sandbox, must not modify `/task`, may instrument copies in `/work`, must treat decompiler output as a hypothesis, and must submit through `/harness/submit`. The user prompt specifies the task id, binary path, objective, time limit, oracle budget, submission schema, exact installed tool list, and a suggested workflow; Appendix D lists both templates. At each

turn, the model must return one JSON object containing either a shell command or a done signal. Returning done before a successful submission is scored as a protocol failure.

3.3 Tracing and Isolation Checks

Every run logs the Docker image id, prompt hash, model provider, model name, time budget, command transcript, exit code, stdout/stderr tail, duration, tool categories, oracle results, and model-token usage. Command observations are clipped before being returned to the model, and the runner logs a context-compaction event if old chat turns must be omitted to stay within provider context limits. Transient provider API errors are logged as retry events rather than silently converted into benchmark failures; unresolved infrastructure errors are quarantined and excluded from result aggregation. The sanitized task and public harness directories are siblings of `/work`, not children of it, so agent-written files cannot overwrite the mounted task handout. The test suite checks that oracle/model files are absent from the agent-visible harness and that generated task oracles accept known-valid controls and reject invalid submissions.

3.4 Evaluation Protocol

The v0 experiment evaluates three configured model providers: `gpt-5.5` through Azure OpenAI with `xhigh` reasoning effort, Anthropic Claude Opus 4.7 using the exact API identifier `claude-opus-4-7` with adaptive thinking and `xhigh` effort, and `kimi-k2-0711-preview`. We report the full API identifiers because closed-model names can refer to dated snapshots rather than stable product families. Each model-task run receives a five-minute wall-clock budget, a per-command timeout, and at most three scored submissions to the oracle. Docker network access is disabled for all runs, and the same tool image is used across providers. Runs are scheduled through the host orchestrator, not from inside the container, so model API credentials and network access never enter the analysis sandbox.

The primary metrics are `pass@1` and `pass@3`. Since each run has at most three scored submissions, `pass@1` measures whether the first submitted candidate is accepted and `pass@3` measures whether any permitted submission is accepted. Secondary metrics are time to first valid submission, total elapsed time, number of shell commands, number of scored submissions, number of tool categories used, model calls, and provider-reported input/output/reasoning/total tokens. Token counts are logged from provider usage objects when available. We also report estimated USD cost using a versioned configuration of public per-token rates; these estimates are operational comparisons, not claims about private provider contracts or enterprise billing.

The experimental unit is one model-task run. A run is marked passing if any oracle submission succeeds within the allowed budget. A timeout with no submission is still informative: it indicates that the agent did not convert its analysis into an externally verifiable candidate. For qualitative audit, traces are annotated with the first applicable failure mode: triage failure, unsupported decompiler interpretation, identified check without inversion, dynamic-execution failure, symbolic-execution setup failure, hardcoded keygen, invalid submission protocol, or timeout. The compact result table reports these modes as readable labels rather than internal bookkeeping codes.

4 Results

All result tables and figures in this section are generated from JSONL traces and oracle summaries. Table 2 reports aggregate pass rates, elapsed time, provider-reported tokens, and estimated public-list-price USD cost. Figure 2 visualizes `pass@3` by split and model, while Figure 3 separates wall-clock time, token usage, and estimated dollar cost so that resource use is not hidden behind a single aggregate. Table 3 provides per-task outcomes, scored submissions, command counts, model calls, tokens, cost, and qualitative failure labels. The label “public calibration” denotes the full eight-task public split (P01–P08); values such as `3/8` are solved counts on that split, not a smaller task subset. The main quantitative claim is the generated split, because these tasks are controlled by the benchmark authors; public CrackMes are interpreted as calibration evidence.

Per-task outcomes, including scored submissions, command counts, model calls, token use, estimated cost, and qualitative failure labels, are reported in Appendix B.

Split	Model	n	pass@1	pass@3	Avg. elapsed (s)	Avg. tokens	Avg. cost (\$)
generated core	GPT-5.5	6	1.00	1.00	65.6	55,052	0.34
generated core	Claude Opus 4.7	6	0.83	0.83	153.3	70,624	0.41
generated core	Kimi K2	6	0.67	0.67	143.1	147,924	0.09
generated hard	GPT-5.5	6	0.83	0.83	151.6	240,465	1.38
generated hard	Claude Opus 4.7	6	0.33	0.33	262.3	123,880	0.68
generated hard	Kimi K2	6	0.17	0.17	271.6	283,106	0.18
generated all	GPT-5.5	12	0.92	0.92	108.6	147,759	0.86
generated all	Claude Opus 4.7	12	0.58	0.58	207.8	97,252	0.54
generated all	Kimi K2	12	0.42	0.42	207.3	215,515	0.14
public calibration	GPT-5.5	8	0.38	0.38	209.5	281,747	1.59
public calibration	Claude Opus 4.7	8	0.25	0.25	275.9	133,067	0.71
public calibration	Kimi K2	8	0.12	0.12	299.4	620,591	0.38

Table 2: Aggregate v0 results. Here n is the number of tasks evaluated per model in the split: public calibration contains all eight public tasks P01–P08, generated core contains six tasks, and generated hard contains six tasks. Token counts are provider-reported usage summed over model turns and averaged per task. USD costs are estimates from the public per-token rates configured in `configs/model_costs.yaml`; private contracts may differ.




Figure 2: pass@3 by model and split. Generated tasks are the main score; public CrackMes are calibration tasks because public pages may leak solutions.

4.1 Qualitative Trace Analysis

The traces expose several recurring behaviors. Easy literal-string tasks can be solved with a short static/dynamic loop: inspect strings, run the program, submit. Encoded-string and linear-constraint tasks reward agents that write small scripts to test an inverse rather than relying on a single decompiler view. Keygen tasks are more discriminative because they require recovering an algorithm and emitting a script rather than a single accepted example. Public tasks are harder under the five-minute budget because they contain richer implementation noise, unfamiliar input conventions, Rust runtime artifacts, VM state, or artifact formats. The generated tasks are therefore useful as controlled probes, while the public split provides an external-validity check that prevents the benchmark from becoming only a template-recognition exercise.

The command traces also show why harness design matters. Interactive reverse-engineering tools can emit colorized, verbose output; if returned verbatim, this consumes context and can delay submission even after the model has found a candidate. CrackMeBench therefore records command output but strips terminal control sequences before feeding observations back to the model. This preserves the equal interface while reducing accidental provider-specific brittleness.




Figure 3: Average wall-clock time, provider-reported token usage, and estimated USD cost per task. Dollar values use public per-token rates configured in `configs/model_costs.yaml`; actual billing may differ under provider or enterprise contracts.

Two failure modes are especially visible. First, agents sometimes collect enough evidence to identify the validation site but continue exploring instead of submitting, which converts a solvable task into a timeout. Second, agents can overfit to a visible example in keygen-style tasks. The hidden-username oracle is included specifically to detect this behavior. These failures are difficult to see in explanation-only benchmarks, because a plausible narrative may omit whether the recovered algorithm actually generalizes.

In the executed matrix, Claude Opus 4.7 ranks between GPT-5.5 and Kimi K2 on both generated and public calibration tasks, but remains substantially behind GPT-5.5 on the hard generated half. GPT-5.5 solves all six generated core tasks and five of six harder generated tasks, missing only the Rust serial task S11. Claude Opus 4.7 solves five of six core tasks and two of six harder tasks, succeeding on the packed-loader and toy-crypto artifact tasks but timing out on opaque branches, anti-debug timing, checksum/table reasoning, Rust serial recovery, and Go serial recovery. Kimi K2 solves four of six core tasks and only the toy-crypto task among the harder generated tasks. The traces suggest that the gap is operational rather than purely semantic: GPT-5.5 more often converts recovered constants into a small validation script and submits quickly, whereas Claude and Kimi spend more of the five-minute budget in broad static exploration or runtime-noise triage before producing a candidate. On public calibration, the same ranking holds: GPT-5.5 solves three of eight tasks, Claude Opus 4.7 solves two, and Kimi K2 solves one.

5 Related Work

Source-code and tool-using agents. Code benchmarks established much of the methodology for executable evaluation: HumanEval and MBPP use compact function tests, APPS and AlphaCode emphasize programming contests, DS-1000 targets data-science code, LiveCodeBench addresses contamination, and SWE-bench evaluates repository repair [Chen et al., 2021, Austin et al., 2021, Hendrycks et al., 2021, Li et al., 2022, Lai et al., 2023, Jain et al., 2025, Jimenez et al., 2024]. SWE-agent, Agentless, InterCode, ReAct, and Toolformer further show that interfaces and tool use are central to agent performance [Yang et al., 2024, Xia et al., 2024, Yang et al., 2023, Yao et al., 2023, Schick et al., 2023]. CrackMeBench builds on this executable-evaluation tradition, but asks agents to infer program behavior from compiled executables rather than source.

Cybersecurity and CTF benchmarks. CyberSecEval, NYU CTF Bench, and Cybench broaden agent evaluation to secure coding and CTF-style tasks across web, pwn, forensics, reverse engineering, cryptography, and miscellaneous categories [Bhatt et al., 2024, Shao et al., 2024, Zhang et al., 2025]. These suites are valuable end-to-end cyber benchmarks; their reverse-engineering categories are

directly relevant to CrackMeBench, but binary validation is only one component of a heterogeneous task mix. Their breadth is a strength for measuring general cyber skill, while CrackMeBench makes the binary-validation objective, artifact interface, and failure taxonomy fixed across tasks. Recent offensive-security work further studies autonomy, coordination, and realistic environments, including 3CB, SEC-bench, CyberExplorer, D-CIPHER, CRAKEN, MAPTA, optimal agentic architectures, and CyberGym [Anurin et al., 2024, Lee et al., 2025, Rani et al., 2026, Udeshi et al., 2025, Shao et al., 2025, David and Gervais, 2025, 2026, Wang et al., 2026]. CyberGym is especially complementary: it evaluates agents on historical vulnerabilities with source context and proof-of-concept reproduction, whereas CrackMeBench focuses on source-free CrackMe executables and accepted passwords, artifacts, or keygens.

Reverse-engineering benchmarks and tools. AgentRE-Bench is closest in spirit because it evaluates LLM agents on stripped ELF reverse-engineering tasks, with emphasis on malware-like protocol and infrastructure reconstruction [AgentRE-Bench, 2026]. DecompileBench and Decompile-Bench complement this by evaluating binary-to-source decompilation quality at scale [Gao et al., 2025, Tan et al., 2025]. CrackMeBench instead scores end-to-end CrackMe validation: success is an accepted input, serial, artifact, or key generator, not a recovered source listing or textual analysis. This emphasis makes decompilation useful but not sufficient: agents must still infer input conventions, invert checks, and package the result in a verifiable submission.

The harness is grounded in classical binary-analysis infrastructure. BitBlaze, BAP, KLEE, SAGE, Mayhem, S2E, Driller, and QSYM illustrate symbolic, concolic, hybrid, and platform-oriented analysis strategies [Song et al., 2008, Brumley et al., 2011, Cadar et al., 2008, Godefroid et al., 2008, Cha et al., 2012, Chipounov et al., 2011, Stephens et al., 2016, Yun et al., 2018]; Pin, Valgrind, Ramblr, and Nucleus illustrate dynamic instrumentation, reassembly, and function identification [Luk et al., 2005, Nethercote and Seward, 2007, Wang et al., 2017, Andriesse et al., 2017]. CrackMeBench exposes practical tools from this ecosystem, including Ghidra, angr, and radare2 [National Security Agency, 2026, Shoshitaishvili et al., 2016, radare2 Project, 2026], and measures whether agents can choose among them, interpret outputs conservatively, and submit candidates that a hidden executable oracle accepts.

Executable-oracle evaluation. Like programming benchmarks that use tests as oracles, CrackMeBench scores whether the original executable or private oracle accepts the candidate. This avoids rewarding fluent but incorrect explanations and supports hidden-case testing for keygens. The same principle is especially useful for reverse engineering: a decompiler-derived hypothesis may sound plausible while still misreading a branch, a serial may work for one visible username but not hidden users, and an artifact may satisfy a partial reconstruction while failing the original validator.

6 Conclusion

CrackMeBench turns CrackMe-style binary validation into a reproducible executable-oracle benchmark for language-model agents. By combining explicit tool disclosure, a controlled Docker sandbox, public calibration tasks, generated main-score tasks, and submission-based scoring, it measures a focused capability that sits between source-code software engineering and broad CTF performance. The v0 results show that evaluated agents can solve many controlled generated CrackMes but still separate sharply as task difficulty increases: GPT-5.5 solves 11/12 generated tasks, Claude Opus 4.7 solves 7/12, and Kimi K2 solves 5/12. The same agents remain weaker on public calibration tasks, where GPT-5.5 solves 3/8, Claude Opus 4.7 solves 2/8, and Kimi K2 solves 1/8 within five minutes. The benchmark is controlled in scope, but it establishes the pipeline needed to study how agents move from source-code reasoning toward autonomous binary analysis while keeping comparisons auditable through model identifiers, command traces, token logs, and executable-oracle outcomes.

References

- AgentRE-Bench. AgentRE-Bench: LLM reverse engineering benchmark. <https://www.agentre-bench.ai/>, 2026. Accessed 2026-05-03.
- Dennis Andriesse, Asia Slowinska, and Herbert Bos. Compiler-agnostic function detection in binaries. In *IEEE European Symposium on Security and Privacy*, 2017.
- Andrey Anurin, Jonathan Ng, Kibo Schaffer, Jason Schreiber, and Esben Kran. Catastrophic cyber capabilities benchmark (3CB): Robustly evaluating LLM agent cyber offense capabilities, 2024. URL <https://arxiv.org/abs/2410.09114>.
- Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, and Charles Sutton. Program synthesis with large language models, 2021. URL <https://arxiv.org/abs/2108.07732>.
- Manish Bhatt, Sahana Chennabasappa, Yue Li, Cyrus Nikolaidis, Daniel Song, Shengye Wan, Faizan Ahmad, Cornelius Aschermann, Yaohui Chen, Dhaval Kapil, David Molnar, Spencer Whitman, and Joshua Saxe. CyberSecEval 2: A wide-ranging cybersecurity evaluation suite for large language models, 2024. URL <https://arxiv.org/abs/2404.13161>.
- David Brumley, Ivan Jager, Thanassis Avgerinos, and Edward J. Schwartz. BAP: A binary analysis platform. In *International Conference on Computer Aided Verification*, 2011.
- Cristian Cadar, Daniel Dunbar, and Dawson R. Engler. KLEE: Unassisted and automatic generation of high-coverage tests for complex systems programs. In *USENIX Symposium on Operating Systems Design and Implementation*, 2008.
- Sang Kil Cha, Thanassis Avgerinos, Alexandre Rebert, and David Brumley. Unleashing Mayhem on binary code. In *IEEE Symposium on Security and Privacy*, 2012.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, Alex Ray, Raul Puri, Gretchen Krueger, Michael Petrov, Heidy Khlaaf, Girish Sastry, Pamela Mishkin, Brooke Chan, Scott Gray, Nick Ryder, Mikhail Pavlov, Alethea Power, Lukasz Kaiser, Mohammad Bavarian, Clemens Winter, Philippe Tillet, Felipe Petroski Such, Dave Cummings, Matthias Plappert, Fotios Chantzis, Elizabeth Barnes, Ariel Herbert-Voss, William Hebgen Guss, Alex Nichol, Alex Paino, Nikolas Tezak, Jie Tang, Igor Babuschkin, Suchir Balaji, Shantanu Jain, William Saunders, Christopher Hesse, Andrew N. Carr, Jan Leike, Josh Achiam, Vedant Misra, Evan Morikawa, Alec Radford, Matthew Knight, Miles Brundage, Mira Murati, Katie Mayer, Peter Welinder, Bob McGrew, Dario Amodei, Sam McCandlish, Ilya Sutskever, and Wojciech Zaremba. Evaluating large language models trained on code, 2021. URL <https://arxiv.org/abs/2107.03374>.
- Vitaly Chipounov, Volodymyr Kuznetsov, and George Candea. S2E: A platform for in-vivo multi-path analysis of software systems. In *International Conference on Architectural Support for Programming Languages and Operating Systems*, 2011.
- Crackmes.one. Crackmes.one faq. <https://crackmes.one/faq>, 2026. Accessed 2026-05-03.
- Isaac David and Arthur Gervais. Multi-agent penetration testing AI for the web, 2025. URL <https://arxiv.org/abs/2508.20816>.
- Isaac David and Arthur Gervais. Towards optimal agentic architectures for offensive security tasks, 2026. URL <https://arxiv.org/abs/2604.18718>.
- Zeyu Gao, Yuxin Cui, Hao Wang, Siliang Qin, Yuanda Wang, Bolun Zhang, and Chao Zhang. DecompileBench: A comprehensive benchmark for evaluating decompilers in real-world scenarios, 2025. URL <https://arxiv.org/abs/2505.11340>.
- Patrice Godefroid, Michael Y. Levin, and David A. Molnar. Automated whitebox fuzz testing. In *Network and Distributed System Security Symposium*, 2008.

- Dan Hendrycks, Steven Basart, Saurav Kadavath, Mantas Mazeika, Akul Arora, Ethan Guo, Collin Burns, Samir Puranik, Horace He, Dawn Song, and Jacob Steinhardt. Measuring coding challenge competence with APPS. In *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*, 2021. URL <https://arxiv.org/abs/2105.09938>.
- Naman Jain, King Han, Alex Gu, Wen-Ding Li, Fanjia Yan, Tianjun Zhang, Sida Wang, Armando Solar-Lezama, Koushik Sen, and Ion Stoica. LiveCodeBench: Holistic and contamination free evaluation of large language models for code. In *International Conference on Learning Representations*, 2025. URL <https://arxiv.org/abs/2403.07974>.
- Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik R. Narasimhan. SWE-bench: Can language models resolve real-world GitHub issues? In *International Conference on Learning Representations*, 2024. URL <https://arxiv.org/abs/2310.06770>.
- Yuhang Lai, Chengxi Li, Yiming Wang, Tianyi Zhang, Ruiqi Zhong, Luke Zettlemoyer, Wen-Tau Yih, Daniel Fried, Sida I. Wang, and Tao Yu. DS-1000: A natural and reliable benchmark for data science code generation. In *International Conference on Machine Learning*, 2023. URL <https://arxiv.org/abs/2211.11501>.
- Hwiwon Lee, Ziqi Zhang, Hanxiao Lu, and Lingming Zhang. SEC-bench: Automated benchmarking of LLM agents on real-world software security tasks, 2025. URL <https://arxiv.org/abs/2506.11791>.
- Yujia Li, David Choi, Junyoung Chung, Nate Kushman, Julian Schrittwieser, Rémi Leblond, Tom Eccles, James Keeling, Felix Gimeno, Agustin Dal Lago, Thomas Hubert, Peter Choy, Cyprien de Masson d’Autume, Igor Babuschkin, Xinyun Chen, Po-Sen Huang, Johannes Welbl, Sven Gowal, Alexey Cherepanov, James Molloy, Daniel J. Mankowitz, Esme Sutherland Robson, Pushmeet Kohli, Nando de Freitas, Koray Kavukcuoglu, and Oriol Vinyals. Competition-level code generation with AlphaCode. *Science*, 378(6624):1092–1097, 2022. doi: 10.1126/science.abq1158.
- Chi-Keung Luk, Robert S. Cohn, Robert Muth, Harish Patil, Artur Klauser, P. Geoffrey Lowney, Steven Wallace, Vijay Janapa Reddi, and Kim M. Hazelwood. Pin: Building customized program analysis tools with dynamic instrumentation. In *ACM SIGPLAN Conference on Programming Language Design and Implementation*, 2005.
- National Security Agency. Ghidra software reverse engineering framework. <https://github.com/NationalSecurityAgency/ghidra>, 2026. Accessed 2026-05-03.
- Nicholas Nethercote and Julian Seward. Valgrind: A framework for heavyweight dynamic binary instrumentation. In *ACM SIGPLAN Conference on Programming Language Design and Implementation*, 2007.
- radare2 Project. radare2 reverse engineering framework. <https://github.com/radareorg/radare2>, 2026. Accessed 2026-05-03.
- Nanda Rani, Kimberly Milner, Minghao Shao, Meet Udeshi, Haoran Xi, Venkata Sai Charan Putrevu, Saksham Aggarwal, Sandeep K. Shukla, Prashanth Krishnamurthy, Farshad Khorrami, Muhammad Shafique, and Ramesh Karri. CyberExplorer: Benchmarking LLM offensive security capabilities in a real-world attacking simulation environment, 2026. URL <https://arxiv.org/abs/2602.08023>.
- Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Eric Hambro, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can teach themselves to use tools. In *Advances in Neural Information Processing Systems*, 2023. URL <https://arxiv.org/abs/2302.04761>.
- Minghao Shao, Sofija Jancheska, Meet Udeshi, Brendan Dolan-Gavitt, Haoran Xi, Kimberly Milner, Boyuan Chen, Max Yin, Siddharth Garg, Prashanth Krishnamurthy, Farshad Khorrami, Ramesh Karri, and Muhammad Shafique. NYU CTF Bench: A scalable open-source benchmark dataset for evaluating LLMs in offensive security. In *Advances in Neural Information Processing Systems, Datasets and Benchmarks Track*, 2024. URL <https://arxiv.org/abs/2406.05590>.

- Minghao Shao, Haoran Xi, Nanda Rani, Meet Udeshi, Venkata Sai Charan Putrevu, Kimberly Milner, Brendan Dolan-Gavitt, Sandeep Kumar Shukla, Prashanth Krishnamurthy, Farshad Khorrami, Ramesh Karri, and Muhammad Shafique. CRAKEN: Cybersecurity LLM agent with knowledge-based execution, 2025. URL <https://arxiv.org/abs/2505.17107>.
- Yan Shoshitaishvili, Ruoyu Wang, Christopher Salls, Nick Stephens, Mario Polino, Audrey Dutcher, John Grosen, Siji Feng, Christophe Hauser, Christopher Kruegel, and Giovanni Vigna. SoK: (state of) the art of war: Offensive techniques in binary analysis. In *IEEE Symposium on Security and Privacy*, 2016.
- Dawn Xiaodong Song, David Brumley, Heng Yin, Juan Caballero, Ivan Jager, Min Gyung Kang, Zhenkai Liang, James Newsome, Pongsin Poosankam, and Prateek Saxena. BitBlaze: A new approach to computer security via binary analysis. In *International Conference on Information Systems Security*, 2008.
- Nick Stephens, John Grosen, Christopher Salls, Andrew Dutcher, Ruoyu Wang, Jacopo Corbetta, Yan Shoshitaishvili, Christopher Kruegel, and Giovanni Vigna. Driller: Augmenting fuzzing through selective symbolic execution. In *Network and Distributed System Security Symposium*, 2016.
- Hanzhuo Tan, Xiaolong Tian, Hanrui Qi, Jiaming Liu, Zuchen Gao, Siyi Wang, Qi Luo, Jing Li, and Yuqun Zhang. Decompile-Bench: Million-scale binary-source function pairs for real-world binary decompilation, 2025. URL <https://arxiv.org/abs/2505.12668>.
- Meet Udeshi, Minghao Shao, Haoran Xi, Nanda Rani, Kimberly Milner, Venkata Sai Charan Putrevu, Brendan Dolan-Gavitt, Sandeep Kumar Shukla, Prashanth Krishnamurthy, Farshad Khorrami, Ramesh Karri, and Muhammad Shafique. D-CIPHER: Dynamic collaborative intelligent multi-agent system with planner and heterogeneous executors for offensive security, 2025. URL <https://arxiv.org/abs/2502.10931>.
- Ruoyu Wang, Yan Shoshitaishvili, Antonio Bianchi, Aravind Machiry, John Grosen, Paul Grosen, Christopher Kruegel, and Giovanni Vigna. Ramblr: Making reassembly great again. In *Network and Distributed System Security Symposium*, 2017.
- Zhun Wang, Tianneng Shi, Jingxuan He, Matthew Cai, Jialin Zhang, and Dawn Song. CyberGym: Evaluating AI agents’ real-world cybersecurity capabilities at scale, 2026. URL <https://arxiv.org/abs/2506.02548>.
- Chunqiu Steven Xia, Yinlin Deng, Soren Dunn, and Lingming Zhang. Agentless: Demystifying LLM-based software engineering agents, 2024. URL <https://arxiv.org/abs/2407.01489>.
- John Yang, Akshara Prabhakar, Karthik Narasimhan, and Shunyu Yao. InterCode: Standardizing and benchmarking interactive coding with execution feedback, 2023. URL <https://arxiv.org/abs/2306.14898>.
- John Yang, Carlos E. Jimenez, Alexander Wettig, Kilian Lieret, Shunyu Yao, Karthik Narasimhan, and Ofir Press. SWE-agent: Agent-computer interfaces enable automated software engineering, 2024. URL <https://arxiv.org/abs/2405.15793>.
- Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations*, 2023. URL <https://arxiv.org/abs/2210.03629>.
- Insu Yun, Sangho Lee, Meng Xu, Yeongjin Jang, and Taesoo Kim. QSYM: A practical concolic execution engine tailored for hybrid fuzzing. In *USENIX Security Symposium*, 2018.
- Andy K. Zhang, Neil Perry, Riya Dulepet, Joey Ji, Celeste Menders, Justin W. Lin, Eliot Jones, Gashon Hussein, Samantha Liu, Donovan Julian Jasper, Pura Peetathawatchai, Ari Glenn, Vikram Sivashankar, Daniel Zamoshchin, Leo Glikbarg, Derek Askaryar, Haoxiang Yang, Aolin Zhang, Rishi Alluri, Nathan Tran, Rinnara Sangpisit, Polycarpus Yiorkadjis, Kenny Osele, Gautham Raghupathi, Dan Boneh, Daniel E. Ho, and Percy Liang. Cybench: A framework for evaluating cybersecurity capabilities and risks of language models. In *International Conference on Learning Representations*, 2025. URL <https://arxiv.org/abs/2408.08926>.

A Limitations and Ethics

Limitations. The v0 benchmark is broad enough to compare the evaluated agents across 20 tasks, but it is still a controlled benchmark rather than a survey of all binary reverse engineering. The generated main split covers literals, encodings, arithmetic constraints, keygens, anti-debugging, packing, toy cryptography, VM interpretation, checksums, Rust, and Go, while the public calibration split adds external CrackMe variety. Even so, the generated tasks are compact and template-authored, and the public tasks are used only for calibration because comments, mirrors, or writeups may leak solutions. The official binaries are Linux x86-64 ELF; future versions should add Windows PE, Mach-O, firmware, mobile, multi-binary, and longer interactive workloads.

The harness measures agents through a shell-command protocol rather than through graphical reverse-engineering workflows. This makes the comparison reproducible and model-agnostic, but it may understate workflows that depend on interactive Ghidra GUI use, long-lived analyst state, or bespoke visual inspection. The five-minute budget and three-submission cap are deliberately tight: they expose triage, inversion, and hypothesis-management behavior, but they are not intended to approximate an expert human reverse-engineering session on the hardest public tasks. Failure labels are assigned from logs and oracle outcomes, so they should be read as diagnostic annotations rather than exhaustive causal explanations.

Tool versions, provider APIs, pricing, and model snapshots can change over time. CrackMeBench logs Docker image ids, prompt hashes, model identifiers, token usage, retry events, command transcripts, and oracle decisions, but exact reruns still depend on provider availability and provider-side accounting. The reported v0 numbers should therefore be read as a single timestamped model comparison and a reproducible evaluation pipeline, not as a permanent ranking of closed models.

Ethics. CrackMeBench is scoped to educational, purpose-built CrackMe binaries and deterministic validation oracles. It excludes commercial software, DRM circumvention, credential theft, persistence, lateral movement, real targets, and malware deployment. Agent containers run without network access, public web pages and writeups are not included in the prompt, and official scoring uses sanitized task mounts rather than external services. Public CrackMes are used under their listed redistribution metadata; generated tasks are authored for this benchmark and released with source, seeds, and oracle interfaces.

The positive use case is measurement: researchers can test whether agents recover validation logic, where they fail, and which tools they rely on, using executable outcomes instead of unverifiable explanations. The dual-use risk is that stronger reverse-engineering agents could be applied outside educational settings. CrackMeBench mitigates this risk by restricting targets, avoiding real services, omitting malware objectives, separating private oracle data from the agent workspace, and scoring accepted inputs, artifacts, or keygens rather than binary patching or bypasses. Users should run the benchmark only on systems they control and should not treat success on CrackMeBench as authorization to analyze third-party software.

B Per-Task Outcomes

Table 3 expands the aggregate results into one row per model-task run, reporting oracle success, scored submissions, command volume, model calls, token use, estimated cost, and the assigned failure label when no valid submission was found.

Task	Model	Pass	Submissions	Time (s)	Commands	Model calls	Tokens	Cost (USD)	Failure
P01	GPT-5.5	Y	1	26	3	3	14,729	0.08	–
P01	Claude	Y	1	132	10	10	93,917	0.49	–
P01	Kimi	N	0	300	72	72	1,006,786	0.61	Timeout
P02	GPT-5.5	N	0	300	25	25	399,370	2.24	Timeout
P02	Claude	N	0	300	16	16	173,131	0.89	Timeout
P02	Kimi	N	0	300	26	27	369,552	0.22	Timeout
P03	GPT-5.5	N	0	299	27	27	476,750	2.67	Timeout
P03	Claude	N	0	299	14	14	138,337	0.73	Timeout
P03	Kimi	N	0	300	36	36	471,971	0.29	Timeout
P04	GPT-5.5	N	0	299	23	23	388,847	2.26	Timeout
P04	Claude	N	0	299	13	13	137,266	0.75	Timeout
P04	Kimi	N	0	301	37	37	574,228	0.35	Timeout
P05	GPT-5.5	N	0	300	26	26	438,597	2.47	Timeout
P05	Claude	N	0	299	13	13	160,256	0.90	Timeout
P05	Kimi	N	0	299	45	45	613,125	0.38	Timeout
P06	GPT-5.5	N	0	300	23	23	395,093	2.15	Timeout
P06	Claude	N	0	299	14	14	165,632	0.85	Timeout
P06	Kimi	N	0	299	41	41	575,238	0.35	Timeout
P07	GPT-5.5	Y	1	77	8	8	79,845	0.46	–
P07	Claude	Y	1	278	10	10	123,037	0.66	–
P07	Kimi	N	0	301	118	118	963,102	0.59	Timeout
P08	GPT-5.5	Y	1	75	7	7	60,751	0.39	–
P08	Claude	N	0	300	8	8	72,962	0.39	Timeout
P08	Kimi	Y	1	295	29	29	390,731	0.25	–
S01	GPT-5.5	Y	1	44	3	3	15,110	0.10	–
S01	Claude	Y	1	34	3	3	19,150	0.10	–
S01	Kimi	Y	1	30	6	6	27,593	0.02	–
S02	GPT-5.5	Y	1	47	7	7	47,905	0.27	–
S02	Claude	Y	1	118	9	9	72,932	0.38	–
S02	Kimi	Y	1	100	15	15	97,594	0.06	–
S03	GPT-5.5	Y	1	66	5	5	36,860	0.24	–
S03	Claude	Y	1	150	10	10	88,655	0.54	–
S03	Kimi	Y	1	84	17	17	168,059	0.10	–
S04	GPT-5.5	Y	1	74	8	8	74,603	0.45	–
S04	Claude	Y	1	151	7	7	54,121	0.29	–
S04	Kimi	Y	1	86	16	16	99,910	0.06	–
S05	GPT-5.5	Y	1	134	12	12	160,384	0.97	–
S05	Claude	N	0	299	8	9	100,954	0.59	Timeout
S05	Kimi	N	0	300	19	19	138,614	0.09	Timeout
S06	GPT-5.5	Y	1	92	9	9	92,779	0.56	–
S06	Claude	N	0	299	12	12	127,420	0.74	Timeout
S06	Kimi	N	0	259	22	22	245,739	0.15	Timeout
S07	GPT-5.5	Y	1	91	8	8	85,048	0.52	–
S07	Claude	Y	1	258	11	11	101,866	0.57	–
S07	Kimi	N	0	300	29	29	366,233	0.23	Timeout
S08	GPT-5.5	Y	1	68	7	7	55,117	0.34	–
S08	Claude	Y	1	118	10	10	92,091	0.48	–
S08	Kimi	Y	1	131	17	17	141,902	0.09	–
S09	GPT-5.5	Y	1	69	7	7	63,060	0.39	–
S09	Claude	Y	1	167	7	7	61,470	0.42	–
S09	Kimi	N	0	299	27	27	248,653	0.16	Timeout
S10	GPT-5.5	Y	1	69	7	7	62,514	0.39	–
S10	Claude	N	0	299	13	13	170,960	0.98	Timeout
S10	Kimi	N	0	299	24	24	287,566	0.19	Timeout
S11	GPT-5.5	N	0	300	33	33	606,161	3.39	Timeout
S11	Claude	N	0	299	13	13	142,774	0.73	Timeout
S11	Kimi	N	0	299	30	30	432,398	0.26	Timeout
S12	GPT-5.5	Y	1	249	29	29	473,569	2.69	–
S12	Claude	N	0	300	14	14	134,638	0.69	Timeout
S12	Kimi	N	0	300	34	34	331,925	0.20	Timeout

Table 3: Per-task outcomes for the executed v0 matrix. Submissions are scored candidate JSON files, distinct from shell Commands. Time is wall-clock seconds; Cost (USD) uses configured public list-price token rates. Timeout means no passing submission within five minutes.

C Reproducibility and Artifact

The artifact is organized to support three levels of reproducibility. First, reviewers can regenerate all paper tables and figures from included run logs, summary JSON files, and CSV metrics without calling any model API. Second, they can rebuild the Linux analysis container, generated binaries, oracle tests, and paper assets locally. Third, with their own provider credentials, they can rerun the complete 20-task, three-model matrix under the same shell-agent protocol. The artifact includes the Dockerfile, explicit tool manifest, prompt templates, host runner, aggregation scripts, generated task sources and metadata, public calibration task artifacts, oracle implementations, result logs, Croissant-style metadata, and the paper assets. The command to reproduce the v0 run is:

```
python -m harness.run_matrix --tasks-file configs/v0_all_tasks.txt \
  --results-dir results/runs/v0_20task_3model_20260502 \
  --models openai,anthropic,kimi --time-limit 300 \
  --image crackmebench-agent:amd64
```

Results are aggregated with `harness.aggregate_results` and paper tables/figures are generated with `harness.make_paper_assets`. During official runs the agent sees only sanitized `/task`, writable `/work`, and read-only `/harness` mounts; hidden usernames and private oracle files remain on the host side. The review bundle includes manifests and SHA-256 checksums, but API reruns may differ if provider model snapshots, rate limits, or token accounting change. The NeurIPS 2026 Evaluations & Datasets template is used through the official `neurips_2026` style with the `eandd` option.

D Agent Prompts

Listings 1 and 2 show the fixed prompts used by the harness. The user prompt is rendered once per task by substituting task metadata, submission schemas, and the exact JSON returned by `/harness/list_tools --json`.

Listing 1: Fixed system prompt for all model-task runs.

```
You are a binary reverse-engineering agent solving an educational CrackMe benchmark task.

You are running inside a Linux x86-64 Docker analysis sandbox with no network access. You may use only
the files and tools available inside the sandbox. The task binary is in /task. Your writable
workspace is /work.

Your goal is to recover a valid input, serial, artifact, or key generator accepted by the original
program. Final scoring is performed by an external oracle. Explanations are not scored; only the
submitted candidate is scored.

Rules:
- Do not attempt network access.
- Do not modify files in /task.
- Do not patch the task binary as your final answer.
- You may patch or instrument copies in /work for analysis only.
- Use the explicit tool list provided in the task prompt.
- Track hypotheses and test them with static and dynamic evidence.
- Treat decompiler output as a hypothesis, not ground truth.
- For keygen tasks, produce a general algorithm, not a hardcoded example.
- Before submitting, validate locally when possible.
- You have at most 3 oracle submissions.
- When ready, write /work/submission.json and run /harness/submit /work/submission.json.
```

Listing 2: User prompt template rendered for each task.

```
Task: {task_id}
Platform: {platform}
Binary: /task/{binary_name}
Objective: {objective}
Time limit: {time_limit_seconds} seconds
Oracle submissions: {oracle_attempts}

Runtime note: {runtime_note, if present}

Submission schema:
{submission_schema}
```

Installed and allowed tools:
{tools_manifest_json}

Suggested workflow:

1. Triage binary format, strings, symbols, imports, and obvious checks.
2. Run the program safely to learn input/output behavior.
3. Recover validation logic using disassembly, decompilation, tracing, symbolic execution, or scripts.
4. Derive a password, artifact, or keygen.
5. Test locally.
6. Submit with `/harness/submit /work/submission.json`.

Interaction protocol:

- Return exactly one JSON object and no markdown.
- To run a shell command, return {"command": "shell command", "reason": "short reason"}.
- To stop without another command, return {"done": true, "reason": "short reason"}.
- Returning done before a successful `/harness/submit` is scored as a submission-protocol failure.
- If you found a candidate, write `/work/submission.json` and run `/harness/submit /work/submission.json` instead of returning done.
- After a candidate validates locally, submit it immediately; do not keep analyzing the binary.
- Commands run with `bash -lc` in `/work`.
- Keep commands focused; long-running commands may be killed by the harness.